

zee5_recsys

July 10, 2024

```
[1]: # Import Libraries
import numpy as np
import pandas as pd
from datetime import datetime
from sklearn.preprocessing import StandardScaler

import warnings
warnings.filterwarnings('ignore')
```

```
[2]: df_movies = pd.read_csv("./zee-movies.dat", delimiter='::',
    ↪encoding="windows-1252")
df_ratings = pd.read_csv("./zee-ratings.dat", delimiter='::')
df_users = pd.read_csv("./zee-users.dat", delimiter='::')
```

```
[3]: df_movies.head()
```

	Movie ID	Title	Genres
0	1	Toy Story (1995)	Animation Children's Comedy
1	2	Jumanji (1995)	Adventure Children's Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama
4	5	Father of the Bride Part II (1995)	Comedy

```
[4]: df_ratings.head()
```

	UserID	MovieID	Rating	Timestamp
0	1	1193	5	978300760
1	1	661	3	978302109
2	1	914	3	978301968
3	1	3408	4	978300275
4	1	2355	5	978824291

```
[5]: df_users.head()
```

	UserID	Gender	Age	Occupation	Zip-code
0	1	F	1	10	48067
1	2	M	56	16	70072
2	3	M	25	15	55117

3	4	M	45	7	02460
4	5	M	25	20	55455

```
[6]: print(df_movies.shape)
      print(df_ratings.shape)
      print(df_users.shape)
```

```
(3883, 3)
(1000209, 4)
(6040, 5)
```

```
[7]: df_movies.nunique()
```

```
[7]: Movie ID      3883
      Title        3883
      Genres        301
      dtype: int64
```

```
[8]: df_users.nunique()
```

```
[8]: UserID        6040
      Gender         2
      Age           7
      Occupation    21
      Zip-code      3439
      dtype: int64
```

```
[9]: print(df_movies.isna().sum())
      print(df_ratings.isna().sum())
      print(df_users.isna().sum())
```

```
Movie ID      0
Title         0
Genres        0
dtype: int64
UserID        0
MovieID       0
Rating        0
Timestamp     0
dtype: int64
UserID        0
Gender        0
Age           0
Occupation    0
Zip-code      0
dtype: int64
```

Filtering only those movies for which the ratings were done most frequently:

```
[10]: df_ratings.MovieID.value_counts().head(1000)
```

```
[10]: MovieID
      2858    3428
      260    2991
      1196    2990
      1210    2883
      480    2672
      ...
      2318    320
      69     319
      2819    319
      1769    319
      1031    319
      Name: count, Length: 1000, dtype: int64
```

```
[11]: select_movies = df_ratings.MovieID.value_counts().head(1000).index.to_list()
      len(select_movies)
```

```
[11]: 1000
```

0.1 Transforming Movie Genres into Binary Features

- Extracting year and title from movie title column
- Let's Split and create OHE Columns from Genre String

```
[12]: import re

def parse_movie_info(movie_string):
    # Regular expression to match the movie name and release year
    pattern = r"^(.*)\s\((\d{4})\)$"
    match = re.match(pattern, movie_string)

    if match:
        movie_name = match.group(1)
        release_year = match.group(2)
        return movie_name, release_year
    else:
        return None, None

# Example usage
input_string = "Inception (2010)"
result = parse_movie_info(input_string)

print(f"Movie Name: {result[0]}")
print(f"Release Year: {result[1]}")
```

Movie Name: Inception

Release Year: 2010

```
[13]: movies = df_movies.copy()
```

```
[14]: movies['movie_name'] = movies['Title'].apply(lambda x: parse_movie_info(x)[0])
      movies['release_year'] = movies['Title'].apply(lambda x: parse_movie_info(x)[1])
      movies.head()
```

```
[14]:
```

	Movie ID	Title	Genres \
0	1	Toy Story (1995)	Animation Children's Comedy
1	2	Jumanji (1995)	Adventure Children's Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama
4	5	Father of the Bride Part II (1995)	Comedy

	movie_name	release_year
0	Toy Story	1995
1	Jumanji	1995
2	Grumpier Old Men	1995
3	Waiting to Exhale	1995
4	Father of the Bride Part II	1995

```
[15]: movies['Genres'] = movies['Genres'].str.split('|')
      movies['Genres']
```

```
[15]:
```

0	[Animation, Children's, Comedy]
1	[Adventure, Children's, Fantasy]
2	[Comedy, Romance]
3	[Comedy, Drama]
4	[Comedy]
...	
3878	[Comedy]
3879	[Drama]
3880	[Drama]
3881	[Drama]
3882	[Drama, Thriller]

Name: Genres, Length: 3883, dtype: object

```
[16]: movies = movies.explode('Genres')
      movies.head()
```

```
[16]:
```

	Movie ID	Title	Genres	movie_name	release_year
0	1	Toy Story (1995)	Animation	Toy Story	1995
0	1	Toy Story (1995)	Children's	Toy Story	1995
0	1	Toy Story (1995)	Comedy	Toy Story	1995
1	2	Jumanji (1995)	Adventure	Jumanji	1995
1	2	Jumanji (1995)	Children's	Jumanji	1995

```
[17]: movies = movies.pivot(index=['Movie ID', 'movie_name', 'release_year'],
    ↪columns='Genres', values='Title')
movies.head()
```

```
[17]: Genres                                Action      Adventure \
Movie ID movie_name      release_year
1      Toy Story      1995      NaN      NaN
2      Jumanji      1995      NaN      Jumanji (1995)
3      Grumpier Old Men      1995      NaN      NaN
4      Waiting to Exhale      1995      NaN      NaN
5      Father of the Bride Part II      1995      NaN      NaN

Genres                                Animation \
Movie ID movie_name      release_year
1      Toy Story      1995      Toy Story (1995)
2      Jumanji      1995      NaN
3      Grumpier Old Men      1995      NaN
4      Waiting to Exhale      1995      NaN
5      Father of the Bride Part II      1995      NaN

Genres                                Children's \
Movie ID movie_name      release_year
1      Toy Story      1995      Toy Story (1995)
2      Jumanji      1995      Jumanji (1995)
3      Grumpier Old Men      1995      NaN
4      Waiting to Exhale      1995      NaN
5      Father of the Bride Part II      1995      NaN

Genres
Comedy \
Movie ID movie_name      release_year
1      Toy Story      1995      Toy Story
(1995)
2      Jumanji      1995
NaN
3      Grumpier Old Men      1995      Grumpier Old Men
(1995)
4      Waiting to Exhale      1995      Waiting to Exhale
(1995)
5      Father of the Bride Part II      1995      Father of the Bride Part II
(1995)

Genres                                Crime Documentary \
Movie ID movie_name      release_year
1      Toy Story      1995      NaN      NaN
2      Jumanji      1995      NaN      NaN
3      Grumpier Old Men      1995      NaN      NaN
```

4	Waiting to Exhale	1995	NaN	NaN
5	Father of the Bride Part II	1995	NaN	NaN

Genres Drama \

Movie ID	movie_name	release_year		
1	Toy Story	1995		NaN
2	Jumanji	1995		NaN
3	Grumpier Old Men	1995		NaN
4	Waiting to Exhale	1995	Waiting to Exhale (1995)	
5	Father of the Bride Part II	1995		NaN

Genres Fantasy Film-Noir \

Movie ID	movie_name	release_year		
1	Toy Story	1995	NaN	NaN
2	Jumanji	1995	Jumanji (1995)	NaN
3	Grumpier Old Men	1995	NaN	NaN
4	Waiting to Exhale	1995	NaN	NaN
5	Father of the Bride Part II	1995	NaN	NaN

Genres Horror Musical Mystery \

Movie ID	movie_name	release_year			
1	Toy Story	1995	NaN	NaN	NaN
2	Jumanji	1995	NaN	NaN	NaN
3	Grumpier Old Men	1995	NaN	NaN	NaN
4	Waiting to Exhale	1995	NaN	NaN	NaN
5	Father of the Bride Part II	1995	NaN	NaN	NaN

Genres Romance \

Movie ID	movie_name	release_year		
1	Toy Story	1995		NaN
2	Jumanji	1995		NaN
3	Grumpier Old Men	1995	Grumpier Old Men (1995)	
4	Waiting to Exhale	1995		NaN
5	Father of the Bride Part II	1995		NaN

Genres Sci-Fi Thriller War Western

Movie ID	movie_name	release_year				
1	Toy Story	1995	NaN	NaN	NaN	NaN
2	Jumanji	1995	NaN	NaN	NaN	NaN
3	Grumpier Old Men	1995	NaN	NaN	NaN	NaN
4	Waiting to Exhale	1995	NaN	NaN	NaN	NaN
5	Father of the Bride Part II	1995	NaN	NaN	NaN	NaN

```
[18]: movies = ~movies.isna()
      movies = movies.astype(int)
      movies
```

[18]: Genres

Movie ID	movie_name	release_year	Action	Adventure	\
1	Toy Story	1995	0	0	
2	Jumanji	1995	0	1	
3	Grumpier Old Men	1995	0	0	
4	Waiting to Exhale	1995	0	0	
5	Father of the Bride Part II	1995	0	0	
...			
3948	Meet the Parents	2000	0	0	
3949	Requiem for a Dream	2000	0	0	
3950	Tigerland	2000	0	0	
3951	Two Family House	2000	0	0	
3952	Contender, The	2000	0	0	

Movie ID	movie_name	release_year	Animation	Children's	\
1	Toy Story	1995	1	1	
2	Jumanji	1995	0	1	
3	Grumpier Old Men	1995	0	0	
4	Waiting to Exhale	1995	0	0	
5	Father of the Bride Part II	1995	0	0	
...			
3948	Meet the Parents	2000	0	0	
3949	Requiem for a Dream	2000	0	0	
3950	Tigerland	2000	0	0	
3951	Two Family House	2000	0	0	
3952	Contender, The	2000	0	0	

Movie ID	movie_name	release_year	Comedy	Crime	Documentary	\
1	Toy Story	1995	1	0	0	
2	Jumanji	1995	0	0	0	
3	Grumpier Old Men	1995	1	0	0	
4	Waiting to Exhale	1995	1	0	0	
5	Father of the Bride Part II	1995	1	0	0	
...			
3948	Meet the Parents	2000	1	0	0	
3949	Requiem for a Dream	2000	0	0	0	
3950	Tigerland	2000	0	0	0	
3951	Two Family House	2000	0	0	0	
3952	Contender, The	2000	0	0	0	

Movie ID	movie_name	release_year	Drama	Fantasy	Film-Noir	\
1	Toy Story	1995	0	0	0	
2	Jumanji	1995	0	1	0	
3	Grumpier Old Men	1995	0	0	0	

4	Waiting to Exhale	1995	1	0	0
5	Father of the Bride Part II	1995	0	0	0
...		
3948	Meet the Parents	2000	0	0	0
3949	Requiem for a Dream	2000	1	0	0
3950	Tigerland	2000	1	0	0
3951	Two Family House	2000	1	0	0
3952	Contender, The	2000	1	0	0

Genres			Horror	Musical	Mystery \
Movie ID	movie_name	release_year			
1	Toy Story	1995	0	0	0
2	Jumanji	1995	0	0	0
3	Grumpier Old Men	1995	0	0	0
4	Waiting to Exhale	1995	0	0	0
5	Father of the Bride Part II	1995	0	0	0
...		
3948	Meet the Parents	2000	0	0	0
3949	Requiem for a Dream	2000	0	0	0
3950	Tigerland	2000	0	0	0
3951	Two Family House	2000	0	0	0
3952	Contender, The	2000	0	0	0

Genres			Romance	Sci-Fi	Thriller \
Movie ID	movie_name	release_year			
1	Toy Story	1995	0	0	0
2	Jumanji	1995	0	0	0
3	Grumpier Old Men	1995	1	0	0
4	Waiting to Exhale	1995	0	0	0
5	Father of the Bride Part II	1995	0	0	0
...		
3948	Meet the Parents	2000	0	0	0
3949	Requiem for a Dream	2000	0	0	0
3950	Tigerland	2000	0	0	0
3951	Two Family House	2000	0	0	0
3952	Contender, The	2000	0	0	1

Genres			War	Western
Movie ID	movie_name	release_year		
1	Toy Story	1995	0	0
2	Jumanji	1995	0	0
3	Grumpier Old Men	1995	0	0
4	Waiting to Exhale	1995	0	0
5	Father of the Bride Part II	1995	0	0
...		
3948	Meet the Parents	2000	0	0
3949	Requiem for a Dream	2000	0	0

3950	Tigerland	2000	0	0
3951	Two Family House	2000	0	0
3952	Contender, The	2000	0	0

[3883 rows x 18 columns]

```
[19]: movies = movies.reset_index()
      movies.head()
```

```
[19]: Genres  Movie ID          movie_name  release_year  Action  Adventure  \
0          1          Toy Story          1995          0          0
1          2          Jumanji          1995          0          1
2          3  Grumpier Old Men          1995          0          0
3          4  Waiting to Exhale          1995          0          0
4          5  Father of the Bride Part II          1995          0          0
```

Genres	Animation	Children's	Comedy	Crime	Documentary	...	Fantasy	\
0	1	1	1	0	0	...	0	
1	0	1	0	0	0	...	1	
2	0	0	1	0	0	...	0	
3	0	0	1	0	0	...	0	
4	0	0	1	0	0	...	0	

Genres	Film-Noir	Horror	Musical	Mystery	Romance	Sci-Fi	Thriller	War	\
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	0	0	1	0	0	0	
3	0	0	0	0	0	0	0	0	
4	0	0	0	0	0	0	0	0	

Genres	Western
0	0
1	0
2	0
3	0
4	0

[5 rows x 21 columns]

```
[20]: movies.set_index('Movie ID', inplace=True, drop='index')
      movies = movies.reset_index()
      movies.head()
```

```
[20]: Genres  Movie ID          movie_name  release_year  Action  Adventure  \
0          1          Toy Story          1995          0          0
1          2          Jumanji          1995          0          1
2          3  Grumpier Old Men          1995          0          0
```

3	4	Waiting to Exhale	1995	0	0
4	5	Father of the Bride Part II	1995	0	0

Genres	Animation	Children's	Comedy	Crime	Documentary	...	Fantasy	\
0	1	1	1	0	0	...	0	
1	0	1	0	0	0	...	1	
2	0	0	1	0	0	...	0	
3	0	0	1	0	0	...	0	
4	0	0	1	0	0	...	0	

Genres	Film-Noir	Horror	Musical	Mystery	Romance	Sci-Fi	Thriller	War	\
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	0	0	1	0	0	0	
3	0	0	0	0	0	0	0	0	
4	0	0	0	0	0	0	0	0	

Genres	Western
0	0
1	0
2	0
3	0
4	0

[5 rows x 21 columns]

0.1.1 Let's Extract Hour from Timestamp

- Extracting the hour from the timestamp can indeed serve as a valuable user feature.
- By incorporating the hour of the day when users provide ratings, it provides a means to capture time-based patterns and behaviors of users.

```
[21]: ratings = df_ratings.copy()
ratings.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000209 entries, 0 to 1000208
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   UserID      1000209 non-null  int64
1   MovieID     1000209 non-null  int64
2   Rating      1000209 non-null  int64
3   Timestamp   1000209 non-null  int64
dtypes: int64(4)
memory usage: 30.5 MB
```

```
[22]: x = 956716541
datetime.fromtimestamp(x).hour
datetime.fromtimestamp(x).weekday()
# r.tail()
```

[22]: 2

```
[23]: ratings['hour'] = ratings['Timestamp'].apply(lambda x: datetime.
    ↪fromtimestamp(x).hour)
ratings['day'] = ratings['Timestamp'].apply(lambda x: datetime.fromtimestamp(x).
    ↪day)
ratings['weekday'] = ratings['Timestamp'].apply(lambda x: datetime.
    ↪fromtimestamp(x).weekday())
ratings['month'] = ratings['Timestamp'].apply(lambda x: datetime.
    ↪fromtimestamp(x).month)
ratings['year'] = ratings['Timestamp'].apply(lambda x: datetime.
    ↪fromtimestamp(x).year)
ratings['date'] = ratings['Timestamp'].apply(lambda x: datetime.
    ↪fromtimestamp(x).date())

ratings.head()
```

```
[23]:   UserID  MovieID  Rating  Timestamp  hour  day  weekday  month  year  \
0        1      1193        5   978300760    3    1         0      1   2001
1        1       661        3   978302109    4    1         0      1   2001
2        1       914        3   978301968    4    1         0      1   2001
3        1      3408        4   978300275    3    1         0      1   2001
4        1      2355        5   978824291    5    7         6      1   2001

      date
0  2001-01-01
1  2001-01-01
2  2001-01-01
3  2001-01-01
4  2001-01-07
```

0.1.2 Now, Let's Read and Merge User Data with Aggregated Values:

```
[24]: users_ratings_stat = df_users.merge(ratings.groupby('UserID').Rating.mean().
    ↪reset_index(), on='UserID')
users_ratings_stat = users_ratings_stat.merge(ratings.groupby('UserID').hour.
    ↪mean().reset_index(), on='UserID')
users_ratings_stat = users_ratings_stat.merge(ratings.groupby('UserID').day.
    ↪mean().reset_index(), on='UserID')
users_ratings_stat = users_ratings_stat.merge(ratings.groupby('UserID').weekday.
    ↪mean().reset_index(), on='UserID')
```

```

users_ratings_stat = users_ratings_stat.merge(ratings.groupby('UserID').month.
↳mean().reset_index(), on='UserID')
users_ratings_stat = users_ratings_stat.merge(ratings.groupby('UserID').year.
↳mean().reset_index(), on='UserID')
users_ratings_stat = users_ratings_stat.merge(ratings.groupby('UserID').
↳Timestamp.mean().reset_index(), on='UserID')
user_rating_count_mapping = ratings.groupby('UserID').Rating.count()
users_ratings_stat['num_ratings'] = users_ratings_stat.UserID.apply(lambda x:
↳user_rating_count_mapping[x])
users_ratings_stat.head()

```

```

[24]:
  UserID  Gender  Age  Occupation  Zip-code  Rating  hour  day \
0      1      F    1          10    48067  4.188679  3.792453  2.471698
1      2      M   56          16    70072  3.713178  2.968992  1.000000
2      3      M   25          15    55117  3.901961  2.215686  1.000000
3      4      M   45           7    02460  4.190476  1.000000  1.000000
4      5      M   25          20    55455  3.146465  11.656566  31.000000

  weekday  month  year  Timestamp  num_ratings
0  1.471698    1.0  2001.0  9.784297e+08         53
1  0.000000    1.0  2001.0  9.782993e+08        129
2  0.000000    1.0  2001.0  9.782978e+08         51
3  0.000000    1.0  2001.0  9.782942e+08         21
4  6.000000   12.0  2000.0  9.782445e+08        198

```

0.1.3 Possible plots:

- Movies release years histogram
- Movies and no. of genres histogram
- Ratings distribution
- Timestamp of rating time graph
- User zipcode mapping
- Occupation histogram for user
- User age histogram
- Gender histogram
- Age to rating correlation
- Age to movie year correlation
- Gender to rating correlation
- Gender to movie year correlation
- Gender + age to rating correlation

```

[25]: # Movies release years histogram
import matplotlib.pyplot as plt
import seaborn as sns

def plot_hist(data, description, figsize = (24, 6)):

```

```

# Create histogram
plt.figure(figsize=figsize)
# plt.hist(movies.release_year, edgecolor='black')
sns.histplot(data, kde=True)

# Add title and labels
plt.title(description)
plt.xlabel('Value')
plt.ylabel('Frequency')
# Rotate x-axis labels
plt.xticks(rotation=90)

# Show the plot
plt.show()


```

```

[26]: movies = movies.dropna()
      movies.release_year = movies.release_year.astype(int)

      # Movies release years histogram
      plot_hist(np.sort(movies.release_year), "Distribution of movies according to_
      ↪their release years")

```



```

[27]: users = df_users.copy()

```

```

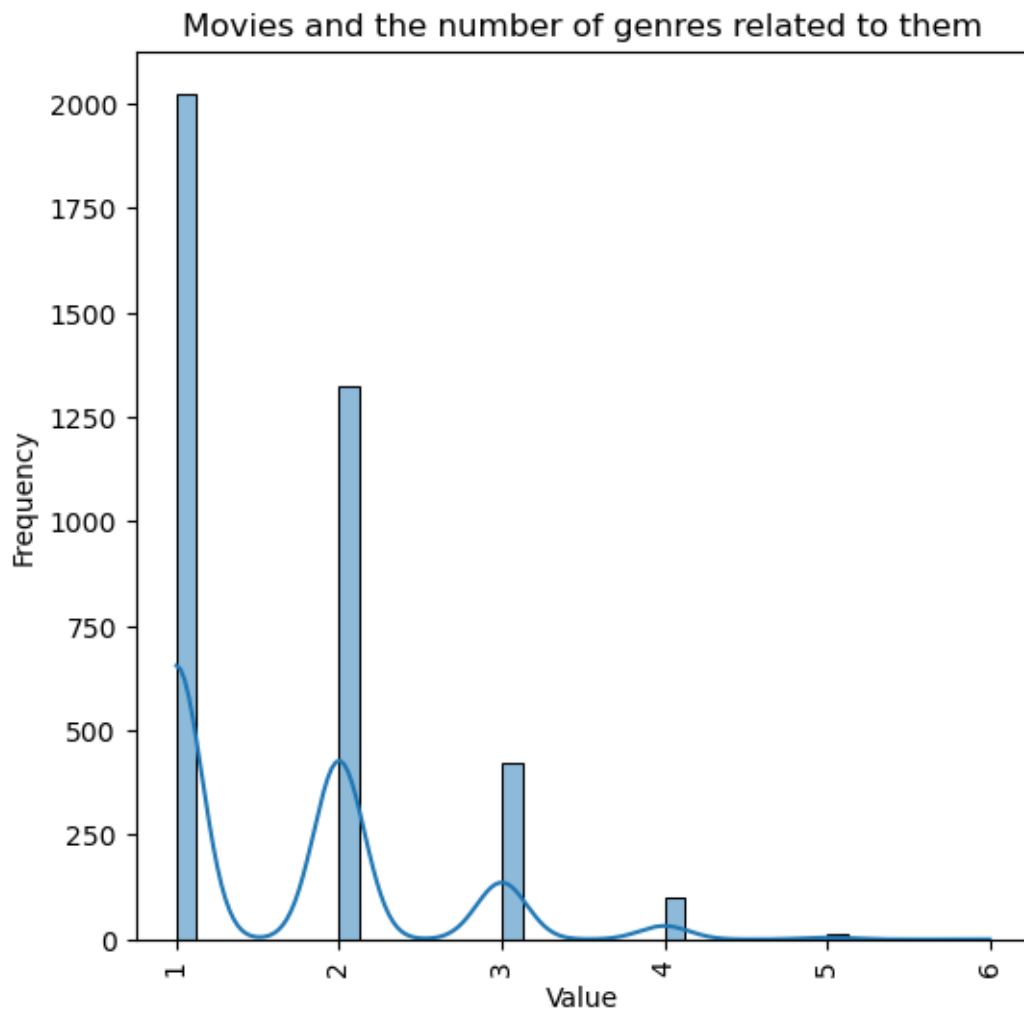
[28]: genre_columns = ['Action', 'Adventure',
                       'Animation', 'Children\'s', 'Comedy', 'Crime', 'Documentary', 'Drama',
                       'Fantasy', 'Film-Noir', 'Horror', 'Musical', 'Mystery', 'Romance',
                       'Sci-Fi', 'Thriller', 'War', 'Western']

def compute_total_genres(row):
    count = 0
    for column in genre_columns:
        count += row[column]
    return count


```

```
movies['genre_count'] = movies.apply(lambda row: compute_total_genres(row),
    ↪axis = 1)
```

```
[29]: # Movies and no. of genres histogram
plot_hist(np.sort(movies.genre_count), "Movies and the number of genres related_
    ↪to them", (6,6))
```




```
[30]: # Ratings distribution
plot_hist(np.sort(ratings.Rating.astype(int)), "Distribution of ratings", (3,3))
```



```
[31]: # Timestamp of rating time graph

import matplotlib.pyplot as plt
num_ratings_vs_date = ratings.groupby('date').Rating.count()

# Plotting the time series
plt.figure(figsize=(10, 6))
num_ratings_vs_date.plot()
plt.title('Number of ratings')
plt.xlabel('Date')
plt.ylabel('Count')
plt.grid(True)
plt.show()
```



```
[32]: users.head()
```

```
[32]:   UserID  Gender  Age  Occupation  Zip-code
0        1      F    1         10    48067
1        2      M   56         16    70072
2        3      M   25         15    55117
3        4      M   45          7    02460
4        5      M   25         20    55455
```

```
[33]: # User zipcode mapping

# import folium
# from geopy.geocoders import Nominatim

# # Initialize geolocator
# geolocator = Nominatim(user_agent="geoapiExercises")

# # Function to get latitude and longitude from zip code
# def get_lat_lon(zip_code):
#     try:
#         location = geolocator.geocode(zip_code)
#         print(location.latitude, location.longitude)
#         return location.latitude, location.longitude
#     except:
```



```

#         print("This is not done")
#         return None, None

# # Apply the function to get coordinates
# users['lat_lon'] = users['Zip-code'].apply(get_lat_lon)
# users[['latitude', 'longitude']] = pd.DataFrame(users['lat_lon'].tolist(),
# ↪ index=users.index)

# # Create a map centered around the mean latitude and longitude
# map_center = [users['latitude'].mean(), users['longitude'].mean()]
# mymap = folium.Map(location=map_center, zoom_start=5)

# # Add markers to the map
# for idx, row in users.iterrows():
#     folium.Marker(location=[row['latitude'], row['longitude']],
# ↪ popup=row['Zip-code']).add_to(mymap)

# # Save the map to an HTML file
# mymap.save("map.html")

```

[34]: # Occupation histogram for user

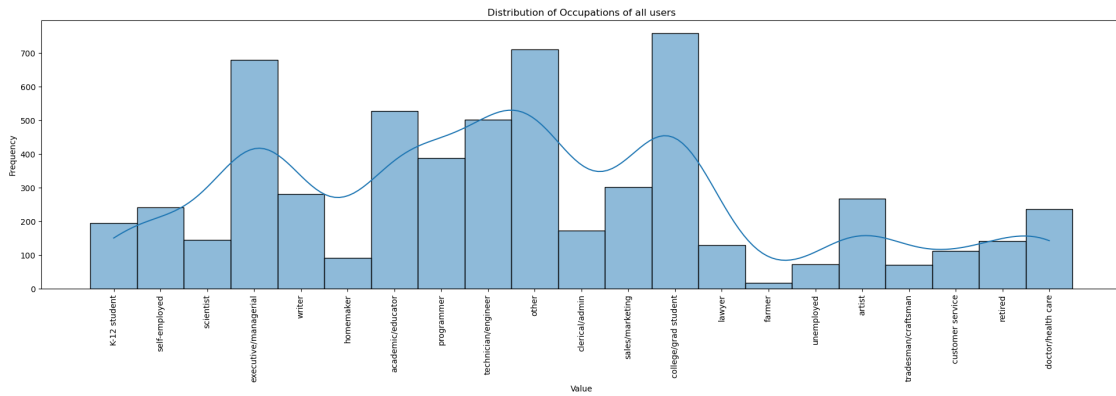
```

occupation_mapping = {
    0: "other",
    1: "academic/educator",
    2: "artist",
    3: "clerical/admin",
    4: "college/grad student",
    5: "customer service",
    6: "doctor/health care",
    7: "executive/managerial",
    8: "farmer",
    9: "homemaker",
    10: "K-12 student",
    11: "lawyer",
    12: "programmer",
    13: "retired",
    14: "sales/marketing",
    15: "scientist",
    16: "self-employed",
    17: "technician/engineer",
    18: "tradesman/craftsman",
    19: "unemployed",
    20: "writer"
}


```

```
users['Occupation_name'] = users.Occupation.apply(lambda x:
↪ occupation_mapping[x])


plot_hist(users.Occupation_name, "Distribution of Occupations of all users")
```



```
[35]: plot_hist(users.Age, "Distribution of ages of all users", figsize=(5,3))
```




```
[36]: plot_hist(users.Gender, "Distribution of genders of all users", figsize=(2,3))
```



```
[37]: # correlation heatmap for user vs ratings info
def gender_to_num(genderName):
    if genderName == 'M':
        return 1
    else:
        return 0
users_ratings_stat['Gender'] = users_ratings_stat.Gender.apply(lambda x:
    ↪gender_to_num(x))
users_ratings_numeric = users_ratings_stat.select_dtypes(include=['int64',
    ↪'float64'])

corr = users_ratings_numeric.corr()

# Create the heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(corr, annot=True, fmt=".2f", cmap='coolwarm', cbar=True)
plt.title('Correlation Heatmap for users rating information')
plt.show()
```



We can see that there is some correlation between Age and Rating and also between num_ratings and Rating. We should keep these columns while making our recommender system

0.1.4 Data analysis over

```
[38]: users['num_ratings'] = users_ratings_stat.UserID.apply(lambda x:
    ↪user_rating_count_mapping[x])
users_movie_rating = ratings.merge(users, on='UserID')
users_movie_rating = users_movie_rating.merge(movies, left_on='MovieID',
    ↪right_on='Movie ID', how="inner")
users_movie_rating = users_movie_rating[['UserID', 'MovieID', 'Rating',
    ↪'movie_name', 'date', 'release_year', 'Action', 'Adventure',
    'Animation', 'Children\'s', 'Comedy', 'Crime', 'Documentary', 'Drama',
    'Fantasy', 'Film-Noir', 'Horror', 'Musical', 'Mystery', 'Romance',
    'Sci-Fi', 'Thriller', 'War', 'Western', 'num_ratings']]
users_movie_rating.rename(columns={'date': 'rating_date'}, inplace=True)
users_movie_rating.head()
```

```
[38]:
```

	UserID	MovieID	Rating	movie_name	rating_date	\
0	1	1193	5	One Flew Over the Cuckoo's Nest	2001-01-01	
1	2	1193	5	One Flew Over the Cuckoo's Nest	2001-01-01	
2	12	1193	4	One Flew Over the Cuckoo's Nest	2000-12-31	
3	15	1193	4	One Flew Over the Cuckoo's Nest	2000-12-30	
4	17	1193	5	One Flew Over the Cuckoo's Nest	2000-12-30	

	release_year	Action	Adventure	Animation	Children's	...	Film-Noir	\
0	1975	0	0	0	0	...	0	
1	1975	0	0	0	0	...	0	
2	1975	0	0	0	0	...	0	
3	1975	0	0	0	0	...	0	
4	1975	0	0	0	0	...	0	

	Horror	Musical	Mystery	Romance	Sci-Fi	Thriller	War	Western	\
0	0	0	0	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	
2	0	0	0	0	0	0	0	0	
3	0	0	0	0	0	0	0	0	
4	0	0	0	0	0	0	0	0	

	num_ratings
0	53
1	129
2	23
3	201
4	211

[5 rows x 25 columns]

0.1.5 Recommend using Pearson coefficient

```
[39]: from scipy.stats import pearsonr

def string_to_numeric_array(s):
    return np.array([ord(char) for char in s])

def top_correlated_movies(input_string, movie_list, top_n=5):
    input_numeric_array = string_to_numeric_array(input_string.lower())
    correlations = []

    for movie in movie_list:
        if(len(movie) > 2):
            shorter_movie_length = min(len(input_numeric_array), len(movie))
            movie_numeric_array = string_to_numeric_array(movie.lower())
            correlation, _ = pearsonr(input_numeric_array[0:
↪shorter_movie_length], movie_numeric_array[0:shorter_movie_length])
```

```

        correlations.append((movie, correlation))

    # Sort correlations based on the absolute value of correlation coefficient
    ↪(higher is better)
    correlations.sort(key=lambda x: abs(x[1]), reverse=True)

    # Get the top n movies with highest correlations
    top_movies = correlations[:top_n]

    return top_movies

top_movies=[]
# Example input from user
user_input = input("Enter a movie name: ")
if(len(user_input) < 2):
    print("Input of length less than 2 is not accepted")
else:
    # Find top 5 movies with highest Pearson correlation coefficient
    top_movies = top_correlated_movies(user_input, list(movies.movie_name))

    # Print the results
    print(f"\nTop 5 movies most correlated with '{user_input}':")
    for i, (movie, correlation) in enumerate(top_movies, 1):
        print(f"{i}. {movie} (Correlation: {correlation:.2f})")

```

Enter a movie name: Toy Story

Top 5 movies most correlated with 'Toy Story':

1. Toy Story (Correlation: 1.00)
2. Toy Story 2 (Correlation: 1.00)
3. Kim (Correlation: 1.00)
4. Kika (Correlation: 0.99)
5. Man Facing Southeast (Hombre Mirando al Sudeste) (Correlation: 0.99)

0.1.6 Recommend using Cosine Similarity

```

[43]: from sklearn.metrics.pairwise import cosine_similarity

user_input = input("Enter a movie to check which are the most relevant movies
    ↪with it: ")

temp = movies[['Action', 'Adventure',
               'Animation', 'Children's', 'Comedy', 'Crime', 'Documentary', 'Drama',
               'Fantasy', 'Film-Noir', 'Horror', 'Musical', 'Mystery', 'Romance',
               'Sci-Fi', 'Thriller', 'War', 'Western']]

# Compute the cosine similarity matrix

```

```

movie_similarity_matrix = cosine_similarity(temp)

# Convert the matrix to a DataFrame for better readability
movie_similarity_matrix_df = pd.DataFrame(movie_similarity_matrix, index=movies.
    ↪movie_name, columns=movies.movie_name)
if user_input in (movie_similarity_matrix_df.index):
    results = movie_similarity_matrix_df.loc[user_input].
    ↪sort_values(ascending=False)[0:6]
    if(user_input in results.index):
        results.drop(user_input, inplace=True)

# Print the results
print(f"\nTop 5 movies most correlated with '{user_input}' according to_
    ↪cosine similarity: ")
for i in range(5):
    movie = results.index[i]
    correlation = results[i]
    print(f"{i+1}. {movie} (Correlation: {correlation:.2f})")
else:
    print("Movie not found in database. Please try with some other movie name.")

```

Enter a movie to check which are the most relevant movies with it: Toy Story

Top 5 movies most correlated with 'Toy Story' according to cosine similarity:

1. American Tail: Fievel Goes West, An (Correlation: 1.00)
2. Rugrats Movie, The (Correlation: 1.00)
3. Toy Story 2 (Correlation: 1.00)
4. Chicken Run (Correlation: 1.00)
5. Adventures of Rocky and Bullwinkle, The (Correlation: 1.00)

```
[44]: movie_similarity_matrix_df.head()
```

```

[44]: movie_name      Toy Story   Jumanji   Grumpier Old Men   \
movie_name
Toy Story           1.000000   0.333333           0.408248
Jumanji             0.333333   1.000000           0.000000
Grumpier Old Men    0.408248   0.000000           1.000000
Waiting to Exhale   0.408248   0.000000           0.500000
Father of the Bride Part II  0.577350   0.000000           0.707107

movie_name      Waiting to Exhale   Father of the Bride Part II   \
movie_name
Toy Story           0.408248           0.577350
Jumanji             0.000000           0.000000
Grumpier Old Men    0.500000           0.707107
Waiting to Exhale   1.000000           0.707107

```

Father of the Bride Part II	0.707107	1.000000
-----------------------------	----------	----------

movie_name	Heat	Sabrina	Tom and Huck	Sudden Death	\
movie_name					
Toy Story	0.0	0.408248	0.408248	0.0	
Jumanji	0.0	0.000000	0.816497	0.0	
Grumpier Old Men	0.0	1.000000	0.000000	0.0	
Waiting to Exhale	0.0	0.500000	0.000000	0.0	
Father of the Bride Part II	0.0	0.707107	0.000000	0.0	

movie_name	GoldenEye	...	Bamboozled	Bootmen	\
movie_name		...			
Toy Story	0.000000	...	0.577350	0.408248	
Jumanji	0.333333	...	0.000000	0.000000	
Grumpier Old Men	0.000000	...	0.707107	0.500000	
Waiting to Exhale	0.000000	...	0.707107	1.000000	
Father of the Bride Part II	0.000000	...	1.000000	0.707107	

movie_name	Digimon: The Movie	Get Carter	Get Carter	\
movie_name				
Toy Story	0.666667	0.000000	0.0	
Jumanji	0.666667	0.000000	0.0	
Grumpier Old Men	0.000000	0.000000	0.0	
Waiting to Exhale	0.000000	0.408248	0.0	
Father of the Bride Part II	0.000000	0.000000	0.0	

movie_name	Meet the Parents	Requiem for a Dream	Tigerland	\
movie_name				
Toy Story	0.577350	0.000000	0.000000	
Jumanji	0.000000	0.000000	0.000000	
Grumpier Old Men	0.707107	0.000000	0.000000	
Waiting to Exhale	0.707107	0.707107	0.707107	
Father of the Bride Part II	1.000000	0.000000	0.000000	

movie_name	Two Family House	Contender, The
movie_name		
Toy Story	0.000000	0.0
Jumanji	0.000000	0.0
Grumpier Old Men	0.000000	0.0
Waiting to Exhale	0.707107	0.5
Father of the Bride Part II	0.000000	0.0

[5 rows x 3882 columns]

```
[45]: users['Gender_numeric'] = users.Gender.apply(lambda x: gender_to_num(x))
temp = users[['Gender_numeric', 'Age', 'Occupation', 'num_ratings']]
```



```
# Compute the cosine similarity matrix
user_similarity_matrix = cosine_similarity(temp)

user_similarity_matrix_df = pd.DataFrame(user_similarity_matrix, index=users.
↳ UserID, columns=users.UserID)

user_similarity_matrix_df.head()
```

```
[45]: UserID      1      2      3      4      5      6      7  \
UserID
1      1.000000  0.923727  0.908048  0.453834  0.990679  0.828739  0.668963
2      0.923727  1.000000  0.988445  0.752428  0.960317  0.979433  0.902616
3      0.908048  0.988445  1.000000  0.781114  0.935468  0.975935  0.899476
4      0.453834  0.752428  0.781114  1.000000  0.539125  0.868662  0.952147
5      0.990679  0.960317  0.935468  0.539125  1.000000  0.884476  0.749728

UserID      8      9     10  ...    6031    6032    6033  \
UserID
1      0.982445  0.977570  0.980841  ...  0.932651  0.918728  0.786609
2      0.973160  0.983634  0.942621  ...  0.991113  0.998667  0.959070
3      0.948017  0.970738  0.902372  ...  0.960161  0.979641  0.968550
4      0.580749  0.627759  0.495429  ...  0.693379  0.747841  0.906890
5      0.998525  0.992962  0.994510  ...  0.972371  0.959825  0.843149

UserID    6034    6035    6036    6037    6038    6039    6040
UserID
1      0.666627  0.980902  0.985615  0.963900  0.350974  0.929051  0.984316
2      0.861595  0.943364  0.923944  0.976151  0.680934  0.991846  0.939738
3      0.913300  0.903334  0.883846  0.941006  0.696861  0.961383  0.901206
4      0.934022  0.497281  0.446037  0.604387  0.987887  0.701414  0.485568
5      0.710423  0.994688  0.991825  0.991119  0.451014  0.969895  0.995244

[5 rows x 6040 columns]
```

0.1.7 Recommend using KNN and Cosine Similarity sklearn

```
[46]: dataset = pd.merge(ratings, movies, left_on='MovieID', right_on='Movie ID')
```

```
[47]: dataset = dataset[['UserID', 'movie_name', 'Rating']]
dataset.head()
```

```
[47]:   UserID      movie_name  Rating
0      1  One Flew Over the Cuckoo's Nest    5
1      2  One Flew Over the Cuckoo's Nest    5
2     12  One Flew Over the Cuckoo's Nest    4
3     15  One Flew Over the Cuckoo's Nest    4
4     17  One Flew Over the Cuckoo's Nest    5
```

```
[48]: table = dataset.pivot_table(index='movie_name', columns='UserID',
↳ values='Rating')
```

```
[49]: table = table.fillna(0)
table.head()
```

```
[49]: UserID          1      2      3      4      5      6      7      8      9      \
movie_name
$1,000,000 Duck      0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
'Night Mother      0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
'Til There Was You  0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
'burbs, The        0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
...And Justice for All 0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0

UserID          10      ... 6031  6032  6033  6034  6035  6036  6037  \
movie_name
$1,000,000 Duck      0.0    ...  0.0    0.0    0.0    0.0    0.0    0.0    0.0
'Night Mother      0.0    ...  0.0    0.0    0.0    0.0    0.0    3.0    0.0
'Til There Was You  0.0    ...  0.0    0.0    0.0    0.0    0.0    0.0    0.0
'burbs, The        4.0    ...  0.0    0.0    0.0    0.0    0.0    0.0    0.0
...And Justice for All 0.0    ...  0.0    0.0    0.0    0.0    0.0    0.0    0.0

UserID          6038  6039  6040
movie_name
$1,000,000 Duck      0.0    0.0    0.0
'Night Mother      0.0    0.0    0.0
'Til There Was You  0.0    0.0    0.0
'burbs, The        0.0    0.0    0.0
...And Justice for All 0.0    0.0    0.0


[5 rows x 6040 columns]
```

```
[50]: table.loc['Toy Story', :].values.reshape(1, -1)
```

```
[50]: array([[5., 0., 0., ..., 0., 0., 3.]])
```

```
[51]: dataset.Rating.value_counts().sort_values().plot(kind='barh')
```

```
[51]: <Axes: ylabel='Rating'>
```



```
[52]: table.values
```

```
[52]: array([[0., 0., 0., ..., 0., 0., 0.],  
        [0., 0., 0., ..., 0., 0., 0.],  
        [0., 0., 0., ..., 0., 0., 0.],  
        ...,  
        [0., 0., 0., ..., 0., 0., 0.],  
        [0., 0., 0., ..., 0., 0., 0.],  
        [0., 0., 0., ..., 0., 0., 0.]])
```

```
[53]: from scipy.sparse import csr_matrix  
  
matrix = csr_matrix(table.values)
```

```
[54]: from sklearn.metrics.pairwise import cosine_similarity  
      from sklearn.neighbors import NearestNeighbors  
  
knn = NearestNeighbors(metric= 'cosine', algorithm= 'brute')  
knn.fit(matrix)
```

```
[54]: NearestNeighbors(algorithm='brute', metric='cosine')
```

```
[55]: user_query_index = np.random.choice(table.shape[1])
      user_query_index
```

```
[55]: 212
```

```
[56]: table.index[user_query_index]
```

```
[56]: 'Artemisia'
```

```
[57]: table.iloc[user_query_index, :]
```

```
[57]: UserID
      1      0.0
      2      0.0
      3      0.0
      4      0.0
      5      0.0
      ...
    6036      0.0
    6037      0.0
    6038      0.0
    6039      0.0
    6040      0.0
      Name: Artemisia, Length: 6040, dtype: float64
```

```
[58]: input_data = table.iloc[user_query_index, :].values.reshape(1, -1)
      distances, indices = knn.kneighbors(input_data, n_neighbors = 6)
```

Generating recommendation using KNN for the selected movie

```
[59]: for i in range(0, len(distances.flatten())):
      # the below line will be printed anyway
      # we always start i = 0
      # just getting the variable i ready to print for which book we are
      # generating the recommendation for:
      if i == 0:
          print('Recommendation for {0}:\n'.format(table.index[user_query_index]))
      else:
          print('{0}: {1}, with distance of {2}.'.format(i, table.index[indices.
↪flatten()[i]], distances.flatten()[i]))
```

Recommendation for Artemisia:

```
1: Conceiving Ada, with distance of 0.6776175236255572:
2: Live Flesh, with distance of 0.7217993936103348:
3: Modern Affair, A, with distance of 0.7321808008165558:
4: School of Flesh, The (L' École de la chair), with distance of
0.7606827894347603:
5: Bloody Child, The, with distance of 0.7673789474003823:
```

0.1.8 Recommend using matrix factorization

Using cmfrec library Collective matrix factorisation for recommender systems Documentation: <https://cmfrec.readthedocs.io/en/latest/>

- cmfrec library requires input in the form of dataframe not as sparse matrix.
- It required 3 columns UserId, ItemId, Rating.
- An instance of the CMF model is created with various hyperparameters:
- method="als": Specifies the alternating least squares (ALS) optimization method, commonly used for matrix factorization in recommendation systems.
- k=2: Sets the number of latent factors to 2, determining the dimensionality of the latent factor space.
- lambda_=0.1: Sets the regularization strength to 0.1. Regularization is used to prevent overfitting in the model.
- user_bias=False: Indicates that user bias terms are not included in the model. User bias represents a user's overall rating tendency.
- item_bias=False: Excludes item bias terms in the model. Item bias represents an item's overall rating tendency.
- verbose=False: Suppresses verbose output, controlling whether the model's training progress is displayed.

```
[60]: rm_raw = ratings[['UserID', 'MovieID', 'Rating']].copy()
rm_raw.columns = ['UserId', 'ItemId', 'Rating'] # Lib requires specific column
names
rm_raw.head()
```

```
[60]:
```

	UserId	ItemId	Rating
0	1	1193	5
1	1	661	3
2	1	914	3
3	1	3408	4
4	1	2355	5

```
[63]: !pip install cmfrec
```

Collecting cmfrec

Using cached cmfrec-3.5.1.post10-cp311-cp311-macosx_14_0_arm64.whl

Collecting cython (from cmfrec)

Obtaining dependency information for cython from <https://files.pythonhosted.org/packages/b6/83/b0a63fc7b315edd46821a1a381d18765c1353d201246da44558175cddd56/Cython-3.0.10-py2.py3-none-any.whl.metadata>

Downloading Cython-3.0.10-py2.py3-none-any.whl.metadata (3.2 kB)

Collecting numpy>=1.25 (from cmfrec)

Obtaining dependency information for numpy>=1.25 from <https://files.pythonhosted.org/packages/01/4a/611a907421d8098d5edc8c2b10c3583796ee8da4156f8f7de52c2f4c9d>

```

90/numpy-2.0.0-cp311-cp311-macosx_14_0_arm64.whl.metadata
  Downloading numpy-2.0.0-cp311-cp311-macosx_14_0_arm64.whl.metadata (60 kB)
        60.9/60.9 kB
1.5 MB/s eta 0:00:00a 0:00:01
Requirement already satisfied: scipy in
/Users/vaibhavmotwani/anaconda3/lib/python3.11/site-packages (from cmfrec)
(1.11.1)
Requirement already satisfied: pandas in
/Users/vaibhavmotwani/anaconda3/lib/python3.11/site-packages (from cmfrec)
(2.0.3)
Collecting findblas (from cmfrec)
  Using cached findblas-0.1.26.post1-py3-none-any.whl
Requirement already satisfied: python-dateutil>=2.8.2 in
/Users/vaibhavmotwani/anaconda3/lib/python3.11/site-packages (from
pandas->cmfrec) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in
/Users/vaibhavmotwani/anaconda3/lib/python3.11/site-packages (from
pandas->cmfrec) (2023.3.post1)
Requirement already satisfied: tzdata>=2022.1 in
/Users/vaibhavmotwani/anaconda3/lib/python3.11/site-packages (from
pandas->cmfrec) (2023.3)
Collecting numpy>=1.25 (from cmfrec)
  Obtaining dependency information for numpy>=1.25 from https://files.pythonhost
ed.org/packages/1a/2e/151484f49fd03944c4a3ad9c418ed193cfd02724e138ac8a9505d056c5
82/numpy-1.26.4-cp311-cp311-macosx_11_0_arm64.whl.metadata
  Downloading numpy-1.26.4-cp311-cp311-macosx_11_0_arm64.whl.metadata (114 kB)
        114.8/114.8
kB 3.7 MB/s eta 0:00:00
Requirement already satisfied: six>=1.5 in
/Users/vaibhavmotwani/anaconda3/lib/python3.11/site-packages (from python-
dateutil>=2.8.2->pandas->cmfrec) (1.16.0)
Downloading Cython-3.0.10-py2.py3-none-any.whl (1.2 MB)
        1.2/1.2 MB
5.4 MB/s eta 0:00:0000:0100:01m
Downloading numpy-1.26.4-cp311-cp311-macosx_11_0_arm64.whl (14.0 MB)
        14.0/14.0 MB
3.6 MB/s eta 0:00:0000:0100:01
Installing collected packages: findblas, numpy, cython, cmfrec
  Attempting uninstall: numpy
    Found existing installation: numpy 1.24.3
    Uninstalling numpy-1.24.3:
      Successfully uninstalled numpy-1.24.3

```

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.

tables 3.8.0 requires blosc2~=2.0.0, which is not installed.

gensim 4.3.0 requires FuzzyTM>=0.4.0, which is not installed.

numba 0.57.1 requires numpy<1.25,>=1.21, but you have numpy 1.26.4 which is incompatible.

Successfully installed cmfrec-3.5.1.post10 cython-3.0.10
findblas-0.1.26.post1 numpy-1.26.4

```
[64]: from cmfrec import CMF
import time

model = CMF(method="als", k=2, lambda_=0.1, user_bias=False, item_bias=False,
↳ verbose=False)

start = time.time()
model.fit(rm_raw)
end = time.time()
print(f"Time Elapsed in trainig: {end-start}")
```

Time Elapsed in trainig: 1.2690629959106445

```
[65]: model.A_.shape, model.B_.shape
```

```
[65]: ((6040, 2), (3706, 2))
```

```
[66]: rm_raw.Rating.mean(), model.glob_mean_
```

```
[66]: (3.581564453029317, 3.581564426422119)
```

```
[67]: rm = ratings.pivot(index = 'UserID', columns = 'MovieID', values = 'Rating').
↳ fillna(0)
rm.head()
```

```
[67]: MovieID  1      2      3      4      5      6      7      8      9     10     ...  \
UserID
1          5.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    ...
2          0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    ...
3          0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    ...
4          0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0    ...
5          0.0    0.0    0.0    0.0    0.0    2.0    0.0    0.0    0.0    0.0    ...

MovieID  3943  3944  3945  3946  3947  3948  3949  3950  3951  3952
UserID
```

1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
5	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

[5 rows x 3706 columns]

```
[68]: from sklearn.metrics import mean_squared_error as mse
rm__ = np.dot(model.A_, model.B_.T) + model.glob_mean_
print("MSE for the above model: " + str(mse(rm.values[rm > 0], rm__[rm > 0])**0.
↪5))
```

MSE for the above model: 1.3043536493177685

```
[69]: from sklearn.metrics import mean_absolute_percentage_error as mape
rm__ = np.dot(model.A_, model.B_.T) + model.glob_mean_
print("MAPE for the above model: " + str(mape(rm.values[rm > 0], rm__[rm >
↪0])**0.5))
```

MAPE for the above model: 0.6136484117493751

```
[70]: movies.head()
```

```
[70]: Genres  Movie ID          movie_name  release_year  Action  \
0          1          Toy Story          1995          0
1          2          Jumanji          1995          0
2          3  Grumpier Old Men          1995          0
3          4  Waiting to Exhale          1995          0
4          5  Father of the Bride Part II          1995          0
```

```
Genres  Adventure  Animation  Children's  Comedy  Crime  Documentary  ...  \
0          0          1          1          1          0          0  ...
1          1          0          1          0          0          0  ...
2          0          0          0          1          0          0  ...
3          0          0          0          1          0          0  ...
4          0          0          0          1          0          0  ...
```

```
Genres  Film-Noir  Horror  Musical  Mystery  Romance  Sci-Fi  Thriller  War  \
0          0          0          0          0          0          0          0          0
1          0          0          0          0          0          0          0          0
2          0          0          0          0          1          0          0          0
3          0          0          0          0          0          0          0          0
4          0          0          0          0          0          0          0          0
```

```
Genres  Western  genre_count
0          0          3
1          0          3
```


2	0	2
3	0	2
4	0	1

[5 rows x 22 columns]

```
[71]: # n: no of items to recommend
top_items = model.topN(user=10, n=10)
movies.loc[movies["Movie ID"].isin(top_items)]
```

```
[71]: Genres  Movie ID                                movie_name \
638         643                        Peanuts - Die Bank zahlt alles
883         895                        Venice/Venice
1397        1421                        Grateful Dead
2469        2538                        Dancemaker
2754        2823  Spiders, The (Die Spinnen, 1. Teil: Der Golden...
2842        2911                Grandfather, The (El Abuelo)
3311        3380                        Railroaded!
3462        3531                All the Vermeers in New York
3748        3818                        Pot O' Gold
3822        3892                Anatomy (Anatomie)
```

Genres	release_year	Action	Adventure	Animation	Children's	Comedy	Crime	\
638	1996	0	0	0	0	1	0	
883	1992	0	0	0	0	0	0	
1397	1995	0	0	0	0	0	0	
2469	1998	0	0	0	0	0	0	
2754	1919	1	0	0	0	0	0	
2842	1998	0	0	0	0	0	0	
3311	1947	0	0	0	0	0	0	
3462	1990	0	0	0	0	1	0	
3748	1941	0	0	0	0	1	0	
3822	2000	0	0	0	0	0	0	

Genres	Documentary	...	Film-Noir	Horror	Musical	Mystery	Romance	\
638	0	...	0	0	0	0	0	
883	0	...	0	0	0	0	0	
1397	1	...	0	0	0	0	0	
2469	1	...	0	0	0	0	0	
2754	0	...	0	0	0	0	0	
2842	0	...	0	0	0	0	0	
3311	0	...	1	0	0	0	0	
3462	0	...	0	0	0	0	1	
3748	0	...	0	0	1	0	0	
3822	0	...	0	1	0	0	0	

Genres	Sci-Fi	Thriller	War	Western	genre_count
--------	--------	----------	-----	---------	-------------

638	0	0	0	0	1
883	0	0	0	0	1
1397	0	0	0	0	1
2469	0	0	0	0	1
2754	0	0	0	0	2
2842	0	0	0	0	1
3311	0	0	0	0	1
3462	0	0	0	0	3
3748	0	0	0	0	2
3822	0	0	0	0	1

[10 rows x 22 columns]

Calculating Average OverLap


```
[72]: overlap= []
num_rec = []
n = 20
for user in ratings.UserID.unique():
    top_items = model.topN(user=user, n=n)
    user_movies = ratings.loc[(ratings.UserID==user)].MovieID
    valid_rec = set(top_items).intersection(set(user_movies)) # I can only
    ↪measure by what was in the training data

    _ = len(set(ratings.loc[ratings.UserID==user].sort_values(by='Rating',
    ↪ascending=False).head(n).MovieID).intersection(set(valid_rec)))
    overlap.append(_)
    num_rec.append(len(valid_rec))

print('avg_perc_overlap:', np.array(overlap).mean() / np.array(num_rec).mean())
pd.Series(overlap).hist(bins=20)
```

avg_perc_overlap: 0.3445217280326718

[72]: <Axes: >



K-precision

Out of K predictions, how many of those K were relevant?

```
[73]: overlap=[]


for user in ratings.UserID.unique():
    recommendations = model.topN(user=user, n=100)
    user_movies = ratings.loc[(ratings.UserID==user)].MovieID
    valid_rec = set(recommendations).intersection(set(user_movies)) # I can
    ↪only measure by what was in the training data
    relevant_items = ratings.loc[(ratings.UserID==user) & (ratings.Rating>=4)].
    ↪MovieID
    try:
        _ = len(set(recommendations).intersection(set(relevant_items))) /
    ↪len(valid_rec)
    except:
        _ = 0
    overlap.append(_)

overlap = np.array(overlap)
print('avg:', overlap.mean())
```

```
pd.Series(overlap).hist(bins=20)
```

avg: 0.7941254267523916

[73]: <Axes: >



Item-Item and User-User similarity matrix using MF embeddings with $d = 4$

```
[74]: model = CMF(method="als", k=4, lambda_=0.1, user_bias=False, item_bias=False, verbose=False)
```

```
model.fit(rm_raw)
```

```
print(model.A_.shape, model.B_.shape)
```

(6040, 4) (3706, 4)

```
[75]: embedded_users_array = model.A_  
embedded_users_df = pd.DataFrame(embedded_users_array)  
embedded_users_df['UserID'] = embedded_users_df.index + 1  
embedded_users_df.head()
```

```
[75]:
```

	0	1	2	3	UserID
0	-0.127591	-0.152649	0.073956	0.014985	1
1	-0.166674	-0.115769	-0.262104	-0.270889	2
2	-0.119248	0.011038	-0.332830	0.127275	3
3	0.263413	-0.338235	-0.211824	-0.356052	4
4	0.439278	0.028768	-0.210001	-0.103933	5

```
[76]: embedded_movies_array = model.B_
embedded_movies_df = pd.DataFrame(embedded_movies_array)
embedded_movies_df['MovieID'] = embedded_movies_df.index + 1
embedded_movies_df = embedded_movies_df.merge(movies, left_on='MovieID',
↳right_on='Movie ID', how="inner")
embedded_movies_df = embedded_movies_df[['MovieID', 'movie_name', 0, 1, 2, 3]]
embedded_movies_df.head()
```

```
[76]:
```

	MovieID	movie_name	0	1	2	\
0	1	Toy Story	0.681755	-4.451891	-1.456907	
1	2	Jumanji	0.305415	-1.293005	2.252264	
2	3	Grumpier Old Men	-1.629499	-3.767879	0.057835	
3	4	Waiting to Exhale	-2.297714	-2.023359	-0.242263	
4	5	Father of the Bride Part II	-1.275383	-2.542525	0.067054	

	3
0	0.623703
1	0.792580
2	-1.566576
3	-0.194461
4	0.132657

```
[77]: temp = embedded_movies_df[['0', 1, 2, 3]]

# Compute the cosine similarity matrix
movie_similarity_matrix = cosine_similarity(temp)

# Convert the matrix to a DataFrame for better readability
movie_similarity_matrix_df = pd.DataFrame(movie_similarity_matrix,
↳index=embedded_movies_df.movie_name, columns=embedded_movies_df.movie_name)

movie_similarity_matrix_df.head()
```

```
[77]:
```

	movie_name	Toy Story	Jumanji	Grumpier Old Men	\
	movie_name				
	Toy Story	1.000000	0.243567	0.695979	
	Jumanji	0.243567	1.000000	0.271746	
	Grumpier Old Men	0.695979	0.271746	1.000000	
	Waiting to Exhale	0.522231	0.144460	0.862147	
	Father of the Bride Part II	0.767282	0.405268	0.915132	

movie_name	Waiting to Exhale	Father of the Bride Part II	\
movie_name			
Toy Story	0.522231	0.767282	
Jumanji	0.144460	0.405268	
Grumpier Old Men	0.862147	0.915132	
Waiting to Exhale	1.000000	0.916435	
Father of the Bride Part II	0.916435	1.000000	

movie_name	Heat	Sabrina	Tom and Huck	Sudden Death	\
movie_name					
Toy Story	0.948973	0.807188	0.969814	0.588493	
Jumanji	0.126584	0.124824	0.213427	0.468337	
Grumpier Old Men	0.758269	0.914218	0.798084	0.954136	
Waiting to Exhale	0.724711	0.917738	0.710791	0.734096	
Father of the Bride Part II	0.864234	0.953834	0.882289	0.833508	

movie_name	GoldenEye	...	Predator 2	Running Man, The	\
movie_name					
Toy Story	0.849512	...	0.187683	-0.187683	
Jumanji	0.284281	...	0.787033	-0.787033	
Grumpier Old Men	0.962745	...	0.644434	-0.644434	
Waiting to Exhale	0.752316	...	0.493285	-0.493285	
Father of the Bride Part II	0.893178	...	0.592788	-0.592788	

movie_name	Starman	Brother from Another Planet, The	\
movie_name			
Toy Story	0.068096	-0.342019	
Jumanji	0.824836	0.703748	
Grumpier Old Men	-0.130242	0.183969	
Waiting to Exhale	-0.015271	0.171955	
Father of the Bride Part II	0.183414	0.161865	

movie_name	Alien Nation	Mad Max	\
movie_name			
Toy Story	0.401123	0.942487	
Jumanji	0.551256	0.240298	
Grumpier Old Men	0.815102	0.499973	
Waiting to Exhale	0.507551	0.215608	
Father of the Bride Part II	0.630446	0.532296	

movie_name	Mad Max 2 (a.k.a. The Road Warrior)	\
movie_name		
Toy Story	-0.882572	
Jumanji	-0.051833	
Grumpier Old Men	-0.856117	
Waiting to Exhale	-0.832180	

Father of the Bride Part II

-0.902946

movie_name	Mad Max Beyond Thunderdome	Bird on a Wire \
movie_name		
Toy Story	-0.130158	0.194384
Jumanji	0.393247	-0.520400
Grumpier Old Men	0.291255	0.569964
Waiting to Exhale	-0.097018	0.756488
Father of the Bride Part II	-0.010381	0.487669

movie_name	Angel Heart
movie_name	
Toy Story	0.875937
Jumanji	0.226801
Grumpier Old Men	0.298411
Waiting to Exhale	0.053800
Father of the Bride Part II	0.389199

[5 rows x 3637 columns]

```
[78]: temp = embedded_users_df[[0, 1, 2, 3]]

# Compute the cosine similarity matrix
user_similarity_matrix = cosine_similarity(temp)

# Convert the matrix to a DataFrame for better readability
user_similarity_matrix_df = pd.DataFrame(user_similarity_matrix,
    ↪index=embedded_users_df.UserID, columns=embedded_users_df.UserID)

user_similarity_matrix_df.head()
```

```
[78]: UserID      1      2      3      4      5      6      7  \
UserID
1      1.000000  0.170103 -0.114736 -0.023483 -0.730620  0.198021  0.706207
2      0.170103  1.000000  0.443404  0.576851  0.031150  0.898441 -0.010590
3     -0.114736  0.443404  1.000000 -0.044441  0.024544  0.189157  0.391775
4     -0.023483  0.576851 -0.044441  1.000000  0.630551  0.330476 -0.083172
5     -0.730620  0.031150  0.024544  0.630551  1.000000 -0.211846 -0.415672

UserID      8      9     10  ...    6031    6032    6033  \
UserID
1      0.302334 -0.263075  0.956633 ... -0.087113 -0.199153  0.235181
2     -0.154776  0.357164  0.427887 ...  0.716629  0.190380  0.967050
3      0.455016  0.822323  0.058423 ...  0.701142 -0.132900  0.555875
4      0.082339  0.381510  0.046103 ... -0.090641  0.899136  0.357558
5      0.059542  0.502422 -0.728975 ... -0.226129  0.812116 -0.166764
```

	6034	6035	6036	6037	6038	6039	6040
UserID							
1	-0.490261	-0.385200	-0.309235	-0.260994	0.150330	0.075981	-0.508371
2	-0.398616	0.475633	0.297296	0.676744	0.820793	0.827002	0.249295
3	-0.192374	0.514526	-0.029510	0.297888	-0.071948	0.422620	-0.025062
4	0.440751	0.745867	0.918480	0.917035	0.859293	0.865973	0.847021
5	0.844038	0.784255	0.866427	0.722518	0.276507	0.426876	0.946257

[5 rows x 6040 columns]

Get d=2 embeddings, and plot the results. Write down your analysis from this visualisation. (Compare with other visualization techniques)

```
[80]: # Modelling MF embedding with d = 2
model = CMF(method="als", k=2, lambda_=0.1, user_bias=False, item_bias=False,
↳ verbose=False)
model.fit(rm_raw)
print(model.A_.shape, model.B_.shape)

# Computing embedded users array
embedded_users_array = model.A_
embedded_users_df = pd.DataFrame(embedded_users_array)
embedded_users_df['UserID'] = embedded_users_df.index + 1

# Computing movies users array
embedded_movies_array = model.B_
embedded_movies_df = pd.DataFrame(embedded_movies_array)
embedded_movies_df['MovieID'] = embedded_movies_df.index + 1
embedded_movies_df = embedded_movies_df.merge(movies, left_on='MovieID',
↳ right_on='Movie ID', how="inner")
embedded_movies_df = embedded_movies_df[['MovieID', 'movie_name', 0, 1]]

# Computing cosine similarity for movies array
temp = embedded_movies_df[[0, 1]]
movie_similarity_matrix = cosine_similarity(temp)
movie_similarity_matrix_df = pd.DataFrame(movie_similarity_matrix,
↳ index=embedded_movies_df.movie_name, columns=embedded_movies_df.movie_name)

# Computing cosine similarity for users array
temp = embedded_users_df[[0, 1]]
user_similarity_matrix = cosine_similarity(temp)
user_similarity_matrix_df = pd.DataFrame(user_similarity_matrix,
↳ index=embedded_users_df.UserID, columns=embedded_users_df.UserID)
```


(6040, 2) (3706, 2)

```
[81]: sns.scatterplot(x=embedded_users_df[0], y=embedded_users_df[1])
```



```
# Adding labels and title
plt.xlabel('X axis')
plt.ylabel('Y axis')
plt.title('Scatter Plot of embedded users info')

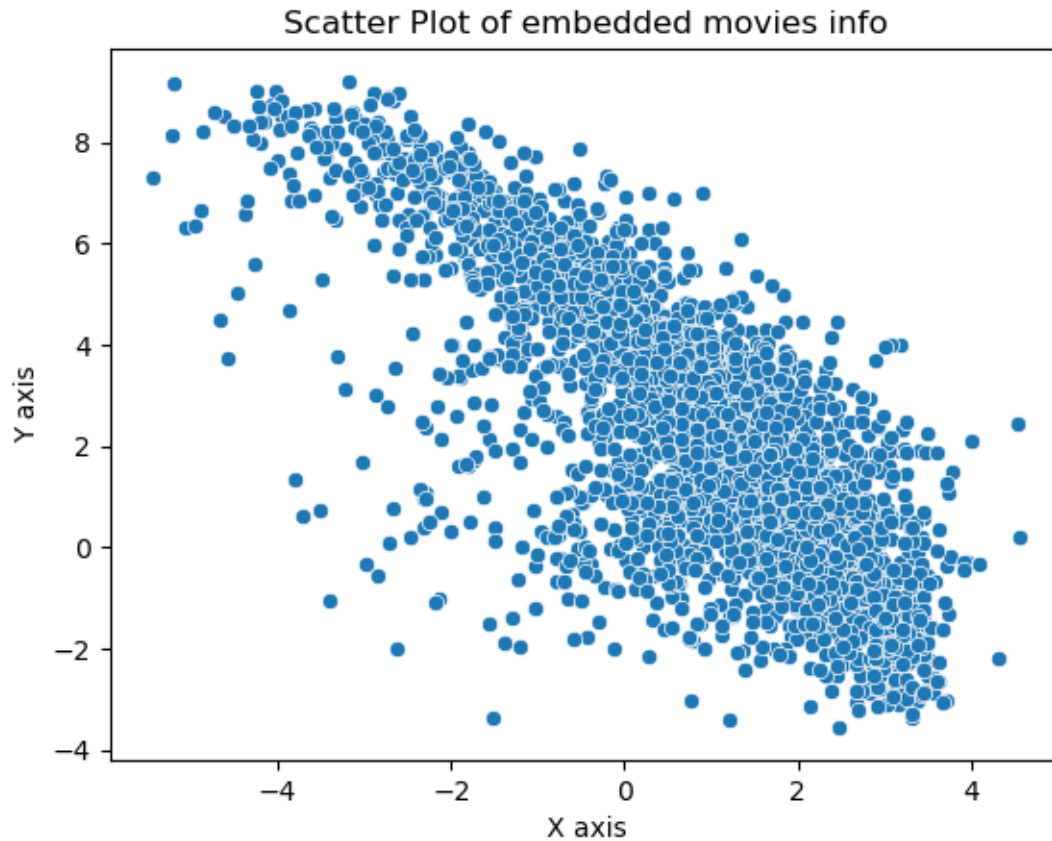
# Show plot
plt.show()
```



```
[82]: sns.scatterplot(x=embedded_movies_df[0], y=embedded_movies_df[1])

# Adding labels and title
plt.xlabel('X axis')
plt.ylabel('Y axis')
plt.title('Scatter Plot of embedded movies info')

# Show plot
plt.show()
```



```
[83]: def get_decade(year):  
    # Ensure the input is an integer  
    try:  
        year = int(year)  
    except ValueError:  
        return "Invalid input. Please enter a valid year."  
  
    # Calculate the decade  
    decade = (year // 10) * 10  
    return decade - 1900  
  
movies['decade'] = movies['release_year'].apply(lambda x: get_decade(x))  
movies.groupby('decade').decade.value_counts().sort_values(ascending=False)
```

```
[83]: decade  
90      2283  
80       597  
70       247  
60       191  
50       168
```

```

100      156
40       126
30       77
20       34
10        3
Name: count, dtype: int64

```

```
[84]: ratings.groupby('MovieID').count().sort_values(by=['UserID'], ascending=False)
movies[movies['Movie ID'] == 2858]
```

```
[84]: Genres  Movie ID      movie_name  release_year  Action  Adventure  Animation \
2789      2858  American Beauty      1999         0         0         0

Genres  Children's  Comedy  Crime  Documentary  ...  Horror  Musical  Mystery \
2789           0        1      0           0  ...      0         0         0

Genres  Romance  Sci-Fi  Thriller  War  Western  genre_count  decade
2789           0      0        0    0      0           2       90

[1 rows x 23 columns]
```

0.2 Questions:

- Users of which age group have watched and rated the most number of movies?
25-34
- Users belonging to which profession have watched and rated the most movies?
College/Student
- Most of the users in our dataset who've rated the movies are Male. (T/F)
True
- Most of the movies present in our dataset were released in which decade? a. 70s b. 90s c. 50s d.80s
90s
- The movie with maximum no. of ratings is _____.
American Beauty
- Name the top 3 movies similar to 'Liar Liar' on the item-based approach.
 1. Castle, The (Correlation: 1.00)
 2. Night Shift (Correlation: 1.00)
 3. EDtv (Correlation: 1.00)
- On the basis of approach, Collaborative Filtering methods can be classified into *-based and* *-based*.
 1. Item

2. User

- Pearson Correlation ranges between ____ to ____ whereas, Cosine Similarity belongs to the interval between ____ to ____.

1. -1 to 1
2. 0 to 1

- Mention the RMSE and MAPE that you got while evaluating the Matrix Factorization model.

MSE for the above MF model: 1.3043536493177685

MAPE for the above MF model: 0.6149695011299233

- Give the sparse 'row' matrix representation for the following dense matrix - $\begin{bmatrix} 1 & 0 \\ 3 & 7 \end{bmatrix}$

Row indices: [0, 1, 1]

Column indices: [0, 0, 1]

Values: [1, 3, 7]

[]: